

2.Convolutional Neural Network

Que 1.Convolutional Neural Networks:-

- A **Convolutional Neural Network (CNN)** is a type of **deep learning model** mainly used for tasks involving images, videos, and other grid-like data.
- CNNs are particularly powerful in **computer vision** because they can automatically detect and learn patterns like edges, textures, shapes, and objects.

Key Components of a Convolutional Neural Network

1 . Convolutional layer

- Convolutional layers use small filters (kernels) that scan across the image.
- These filters help the network detect important features like edges, corners, textures, and shapes.
- The main advantage is that convolution **keeps the spatial relationship** → meaning nearby pixels are processed together, so the network understands **where things are in the image**

Hyper parameters of Convolutional Layer

i) Filter Size (Kernel Size)

- The **size of the filter (kernel)** used to scan the image.
- Common sizes: **3×3, 5×5, 7×7**.
- Smaller filters → capture fine details (edges).
- Larger filters → capture more global features (shapes).

Example: A **3×3 filter** looks at a small patch (like 9 pixels at a time).

ii) Output Depth (Number of Filters)

- The number of filters (kernels) used in the convolutional layer.
- Each filter detects a different feature → **edges, corners, textures, colors**.
- More filters → more **feature maps** → higher **output depth**.

Example: If 32 filters are used, the output depth = **32 feature maps**.

iii) Stride

- The **step size** by which the filter moves/slides across the input.
- Stride = 1 → moves **1 pixel at a time** (output is larger).
- Stride = 2 → moves **2 pixels at a time** (output is smaller).

Example: Stride 1 = detailed output, Stride 2 = faster but less detail.

iv) Zero Padding

- Adding extra **rows/columns of zeros** around the image before convolution.
- Purpose:
 - To **control output size** (same size as input if needed).
 - To avoid losing information at the edges.
- Types:
 - **Valid padding:** No padding → output shrinks.
 - **Same padding:** Adds zeros → output size remains same.

Example: A 5×5 image with 3×3 filter → without padding → output = 3×3 ; with padding → output = 5×5 .

2. Pooling Layers

- Pooling layers make the image **smaller** by reducing its size, which makes the network **faster and easier to train**.
- In **Max Pooling**, we look at a small block of pixels and **keep only the biggest value**, ignoring the rest.
- **Pooling layers** are used to **reduce the size** of the feature maps while keeping the important information.
 - This makes the network:
 - Faster (less computation)
 - Smaller (fewer parameters)
 - More general (less overfitting).

1. Stride in Pooling Layers

- **Stride** = how far the pooling window (like 2×2) moves each time.
- **Stride = 1** → moves **1 step at a time** → bigger output (more detail).
- **Stride = 2** → moves **2 steps at a time** → smaller output (less detail, more compression).

Example (2×2 Max Pool, Stride=2):

Input: 4×4 → Output: 2×2

(It skips more pixels, so the output shrinks).

2. Padding in Pooling Layers

- Sometimes, the input size is not perfectly divisible by the pooling window.
- **Padding** adds extra rows/columns (usually zeros) around the image so that pooling can still cover all areas.
- Padding ensures that no edge pixels are ignored.

Example:

- Input: 5×5 , Pooling window: 2×2
- Without padding → last row/column gets ignored.
- With padding → input becomes 6×6 → pooling works neatly.

There are mainly **Three types of pooling**:

1. Max Pooling

- Selects the **maximum value** from the pooling window.
- Example:

Pooling Window = 2×2

```
csharp Copy Edit

Input:
[1, 3]
[2, 4]

Max Pooling → 4 (largest value)
```

- **Use:** Keeps the strongest feature → widely used in CNNs.

2. Average Pooling

- Takes the **average value** of the pooling window.
- Example:

```
csharp Copy Edit

Input:
[1, 3]
[2, 4]

Average Pooling → (1+3+2+4)/4 = 2.5
```

- **Use:** Smooths out feature maps by considering all values.

3. Min Pooling (less common)

- Selects the **minimum value** from the pooling window.
- Example:

```
csharp Copy Edit

Input:
[1, 3]
[2, 4]

Min Pooling → 1
```

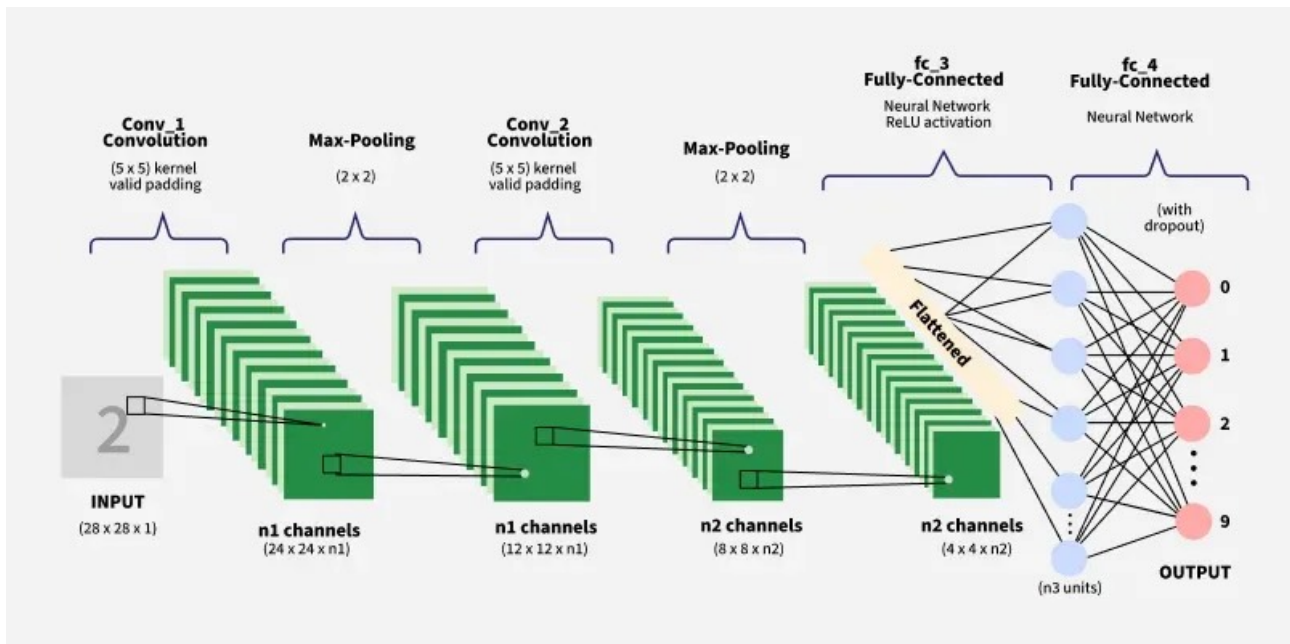
- **Use:** Rare, but sometimes useful for detecting very low-intensity features.
-

3. Activation Function Layer (ReLU)

- They introduce non-linearity to the model by allowing it to learn more complex relationships in the data.

4. Fully Connected (Dense) Layer

- These layers are responsible for making predictions based on the high-level features learned by the previous layers. They connect every neuron in one layer to every neuron in the next layer.



How CNNs Work?

1. **Input Image:** CNN receives an input image which is preprocessed to ensure uniformity in size and format.
2. **Convolutional Layers:** Filters are applied to the input image to extract features like edges, textures and shapes.
3. **Pooling Layers:** The feature maps generated by the convolutional layers are downsampled to reduce dimensionality.
4. **Fully Connected Layers:** The downsampled feature maps are passed through fully connected layers to produce the final output, such as a classification label.
5. **Output:** The CNN outputs a prediction, such as the class of the image.

Functions of Hidden Layers in CNN

Hidden layers are the **intermediate layers** (Convolution + ReLU + Pooling) between input and output. Their main functions are:


1. **Feature Extraction (Convolutional Layers)**
 - Detects patterns like lines, curves, textures, and later more complex objects.
2. **Non-Linearity (Activation Layers)**
 - Adds flexibility by letting the network learn **complex mappings** instead of just linear transformations.
3. **Dimensionality Reduction (Pooling Layers)**
 - Reduces the size of feature maps → makes computation faster and prevents overfitting.
4. **Representation Learning**
 - Each hidden layer learns a **higher-level representation**:
 - Early layers → simple features (edges, colors).
 - Deeper layers → complex features (faces, objects).

Application:-

1. **Image Classification & Recognition** – CNNs classify images into categories like cat, dog, or car.
2. **Object Detection** – They locate and identify multiple objects within an image.
3. **Face Recognition** – CNNs match and verify faces for authentication or tagging.
4. **Medical Image Analysis** – They detect diseases and abnormalities from X-rays, MRIs, and CT scans.
5. **Self-Driving Cars** – CNNs help vehicles recognize lanes, pedestrians, and traffic signals.
6. **Natural Language Processing (NLP)** – They analyze text for sentiment, spam detection, and classification.
7. **Video Analysis** – CNNs identify actions, objects, and events in video footage.
8. **Robotics & Industry** – They guide robots in object detection and quality inspection.
9. **Agriculture** – CNNs detect crop diseases and differentiate between crops and weeds.
10. **Art & Creativity** – They enable style transfer, AI art, and creative image generation.

Que.2 Draw and explain architecture of AlexNet.

- AlexNet is a famous **Convolutional Neural Network (CNN)** that became very popular in 2012 because it won the ImageNet competition with much better accuracy than others.
- It works with large images ($224 \times 224 \times 3$ RGB) and has 5 convolutional layers + 3 fully connected layer



- Structure (8 layers):**
It has **5 convolutional layers** (for learning image features) and **3 fully connected layers** (for making final predictions). Between some layers, it uses **max pooling** to reduce size but keep important features.

input ($224 \times 224 \times 3$) →

- Conv1** → 96 filters (11×11 , stride 4) → ReLU → Max Pooling
- Conv2** → 256 filters (5×5) → ReLU → Max Pooling
- Conv3** → 384 filters (3×3) → ReLU
- Conv4** → 384 filters (3×3) → ReLU
- Conv5** → 256 filters (3×3) → ReLU → Max Pooling
- Fully Connected (FC6)** → 4096 neurons + ReLU + Dropout
- Fully Connected (FC7)** → 4096 neurons + ReLU + Dropout
- Output Layer (FC8)** → 1000 neurons (Softmax for 1000 ImageNet classes)

- Key Tricks in AlexNet:**
 - Uses **ReLU** activation instead of sigmoid/tanh → trains faster.
 - Uses **Dropout** in fully connected layers → prevents overfitting.
 - Uses **Overlapping Pooling** → better feature retention.
 - Trained using **two GPUs in parallel** → faster for large dataset.
- Why it was important:**
Before AlexNet, CNNs were not very popular. AlexNet showed that with the right design and training, CNNs can solve large image recognition problems very well. This started the deep learning revolution in computer vision.
- Comparison:**
AlexNet is simpler (only 8 layers) compared to later models like GoogleNet or ResNet (which have dozens or hundreds of layers). This makes AlexNet easier to train but less powerful.

- **Example (Dog Breed Classification):**

Imagine we want to recognize dog breeds:

- Early layers of AlexNet detect **edges and textures** (like fur patterns, ear shapes).
- Middle layers detect **parts of the face or body**.
- Final fully connected layers combine these features and decide which **dog breed** it

Advantages of AlexNet

1. Breakthrough in Deep Learning
2. High Accuracy
3. Use of ReLU Activation
4. GPU Acceleration
5. Dropout Regularization
6. Overlapping Max Pooling
7. Scalable

Que.3 What is convolution operation ?

- Convolution is a **mathematical operation** where we take an **input (like an image)** and a **filter (kernel)**, then slide the filter over the input to produce an **output (feature map)**.
- Convolution is a **mathematical operation** that combines two functions (or signals) to form a third one.
- In **mathematics/signal processing** → It tells how the shape of one function is modified by another.
- In **CNNs (Convolutional Neural Networks)** → It's the process of sliding a **filter/kernel** over an input image to extract features like edges, corners, textures, etc.

How it Works (Step by Step)

1. Take an **input image** (matrix of numbers, each number = pixel value).
2. Take a **filter (kernel)** (a smaller matrix, e.g., 3×3).
3. Place the filter on a part of the image.
4. Multiply each filter value with the overlapping image pixel values.
5. Add them all → this gives **one number** in the output feature map.
6. Slide the filter across the image and repeat.

Convolution in CNN (Step by Step)

Imagine we have a **5×5 image** (input) and a **3×3 filter (kernel)**:

Step 1: Place filter on top-left of the image

Multiply filter values with the overlapping pixel values.

Step 2: Sum the multiplication results

This sum gives **one pixel** in the output feature map.

Step 3: Slide filter

Move the filter across the image (stride decides step size).

Step 4: Repeat until full image is covered

We get a **smaller output matrix** called the **feature map**.

Explain circular and discrete convolution operation and other type

1) Discrete Convolution

- Used when signals/images are **discrete** (not continuous, but stored in digital form).
- Formula:

$$y[n] = \sum_{k=-\infty}^{\infty} x[k] \cdot h[n - k]$$

where:

- $x[k]$ = input signal (image)
- $h[k]$ = filter (kernel)
- $y[n]$ = output (feature map)

👉 Example:

If input $x = [1, 2, 3]$ and filter $h = [1, 1]$,

$$y = [1, (1 \times 2 + 1 \times 1) = 3, (2 \times 1 + 3 \times 1) = 5, 3]$$

So, $y = [1, 3, 5, 3]$.

2) Circular Convolution

- It is like discrete convolution but assumes the signal is **periodic** (wraps around like a circle).
- Instead of padding with zeros, values from the end of the signal wrap around to the beginning.
- Formula:

$$y[n] = \sum_{k=0}^{N-1} x[k] \cdot h[(n - k) \bmod N]$$

👉 Example:

If $x = [1, 2, 3]$ and $h = [1, 1, 1]$,

- Circular convolution gives result by wrapping indices, so output = $[6, 6, 6]$.

1. Valid Convolution

- No padding is added.
- Output size is **smaller** than input.
- Example: Input (5×5) , Filter $(3 \times 3) \rightarrow$ Output (3×3) .

It only uses **valid overlapping regions** between filter & image.

2. Same Convolution

- Padding is added so that the **output size = input size**.
- Example: Input (5×5) , Filter (3×3) , with padding \rightarrow Output (5×5) .

Keeps the feature map size same as the input.

3. Full Convolution

- Padding is added in such a way that the filter can go **outside the input boundary**.
- Output size is **larger** than input.
- Example: Input (5×5) , Filter $(3 \times 3) \rightarrow$ Output (7×7) .

Used rarely in CNNs, but more common in signal processing.